

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
Lecture: RNN Demo

Lecture Agenda

- Logistics
 - RNNs due **During Finals Time**
- Recurrent Networks
 - *Overview*
 - *Problem Types*
 - *Embeddings*
 - *Types of RNNs*
 - **Demo A**
 - **CNNs and RNNs**
 - **Demo B**
 - **Ethics Case Study**
 - **Course Retrospective**


Last Time

- padding/clipping




Visualization

GloVe produces word vectors with a marked banded structure that is evident upon visualization:



Recurrent Networks

- basic RNN




$$h_t = \tanh(W_A (X_t @ h_{t-1}) + b_A)$$

$$P_t = \text{softmax}(W_P h_t + b_P)$$

<https://arxiv.org/pdf/1603.01521v1.pdf>

Self Test

- T/F: In Recurrent Neural Networks that are “rolled out”, each RNN cell can be run in parallel.
 - A. **True**, state vectors can be added later
 - B. **True**, but parallelization must use forward backward (like Viterbi)
 - C. **False**, state vectors must be found sequentially
 - D. **False**, input changes due to sequential nature of X_t

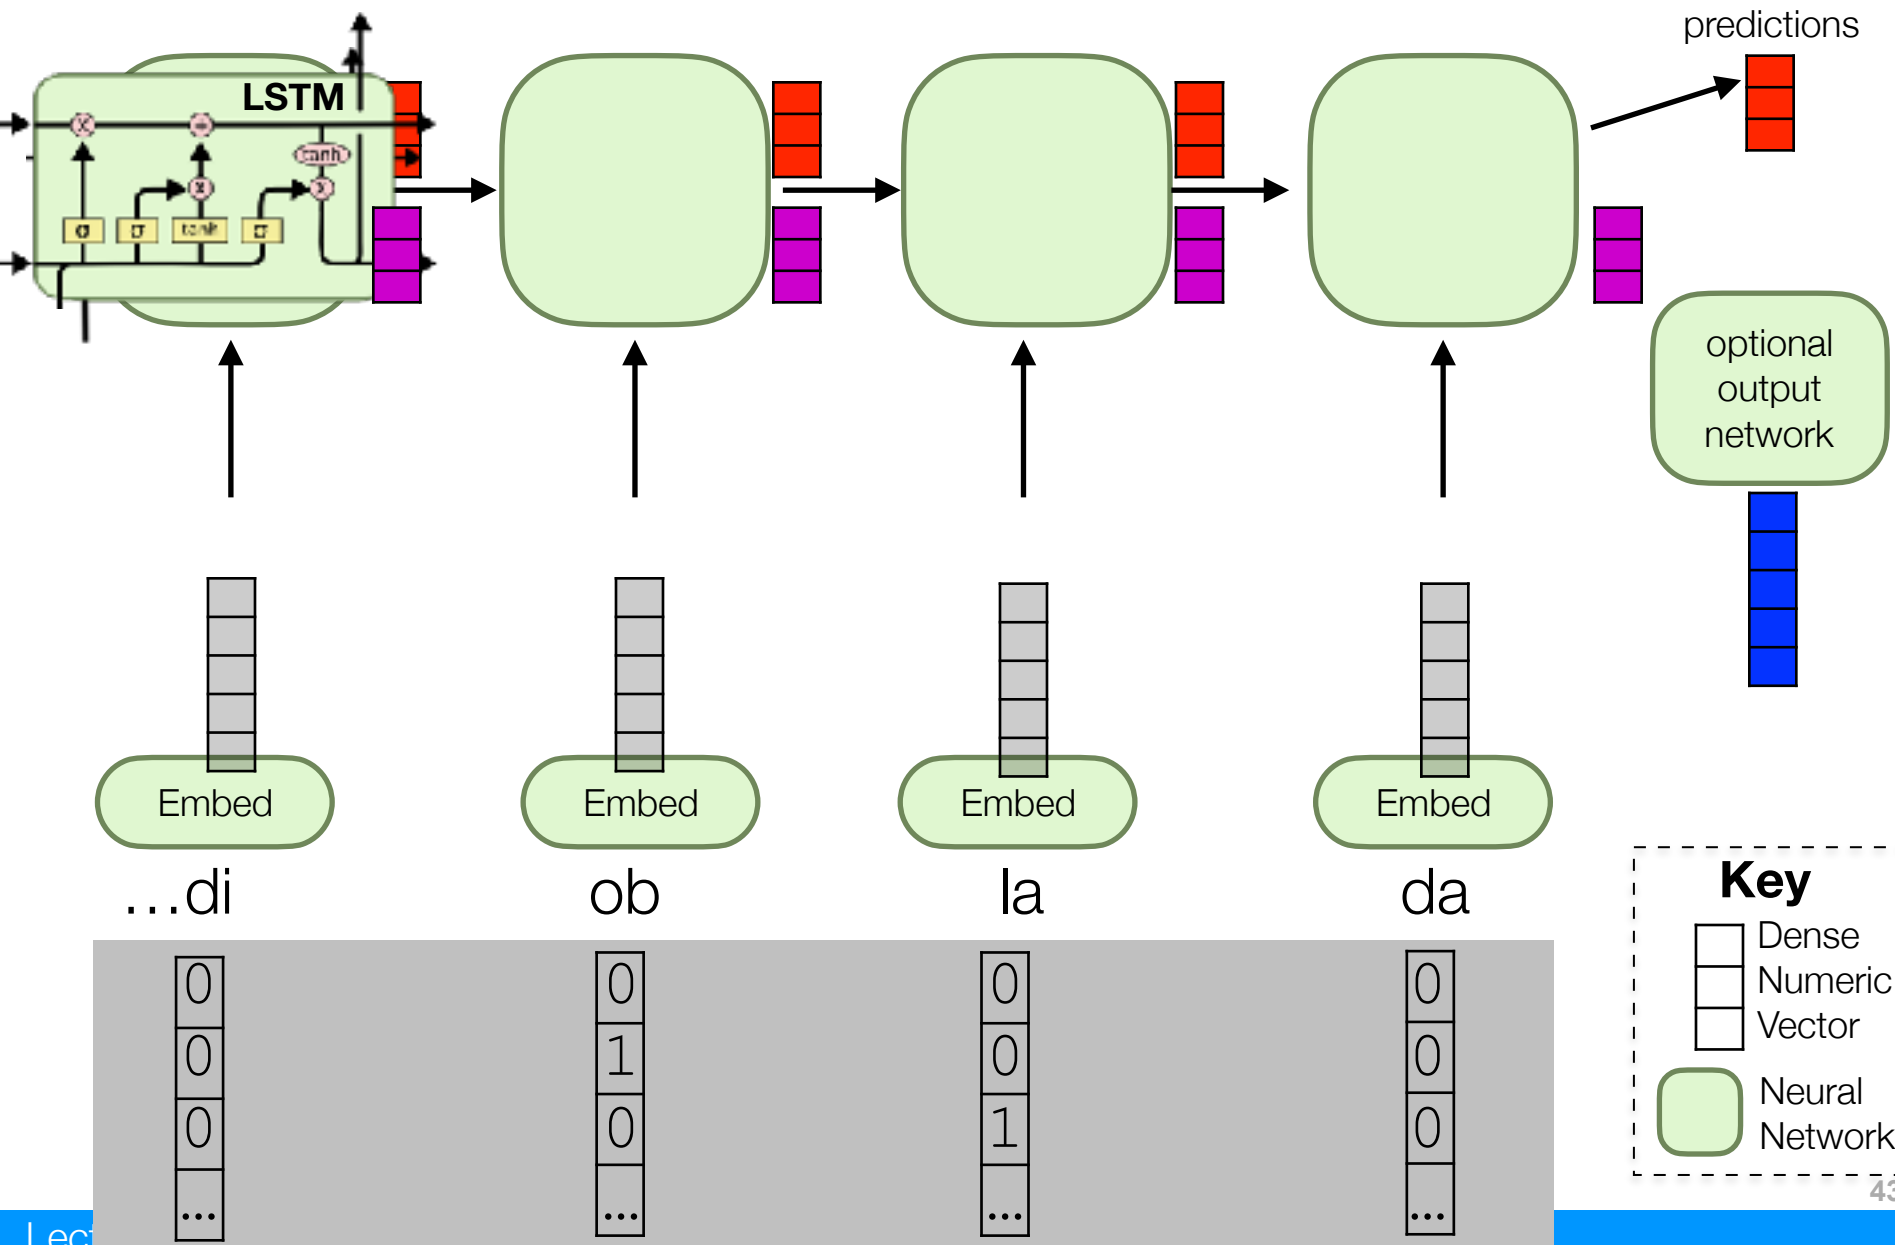


Commonly Used RNN Nodes



I like this version better.

General recurrent flow (many to one)



Recurrent Networks: GRUs

- gated recurrent units

Selectivity controls, gates (**0 or 1**)

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$



past state

current input



selectively remember


with influence

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

remember only past

OR remember with input



σ = sigmoid

\odot = elem. multiplication

Self Test


- What element of the GRU helps with vanishing and exploding gradients?
- A. derivative of σ
- B. no activation function
- C. derivative of ϕ
- D. ϕ

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$


$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$



Derivative of GRU



$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + U x_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

To back propagate, we need sensitivity of state vector, w.r.t previous state

Product Rule

Product Rule

$$\partial s_t / \partial s_{t-1} = (\partial z_t \times s_{t-1}) + (\partial s_{t-1} \times z_t) + \partial \tilde{s}_t - (\partial z_t \times \tilde{s}_t) - (\partial \tilde{s}_t \times z_t)$$

likely vanish

could vanish, depending on ϕ

likely vanish

could vanish, depending on ϕ

hard to vanish unless $z_t = 0$

Recurrent Networks: Gen 1 LSTM

- LSTM prototype

Selectivity controls (**gates, 0 or 1**)

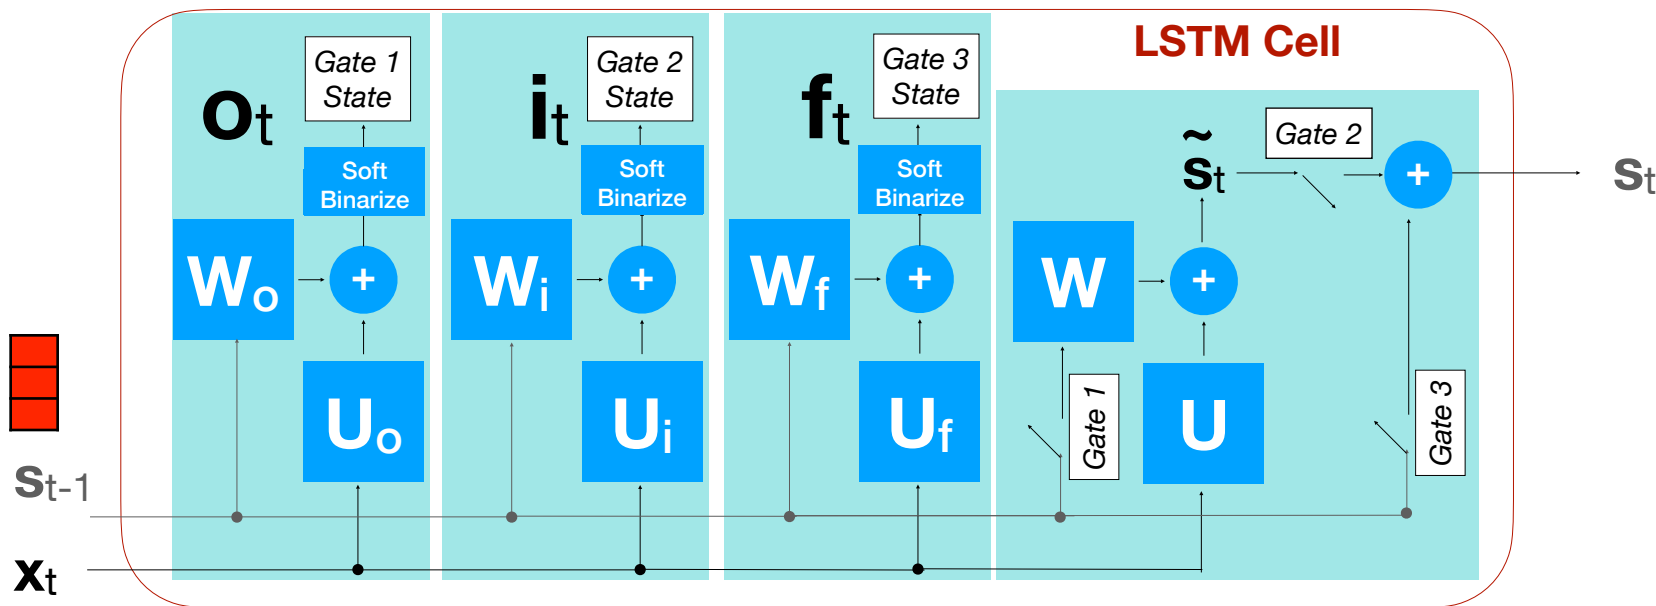
$$o_t = \sigma(W_o s_{t-1} + U_o x_t + b_o)$$

$$i_t = \sigma(W_i s_{t-1} + U_i x_t + b_i)$$

$$f_t = \sigma(W_f s_{t-1} + U_f x_t + b_f)$$

$$\begin{aligned} \tilde{s}_t &= \phi(W(o_t \odot s_{t-1}) + Ux_t + b) \\ s_t &= f_t \odot s_{t-1} + i_t \odot \tilde{s}_t \end{aligned}$$

selectively remember past with influence
selectively remember past with past weighted influence



Recurrent Networks: Gen 2 LSTM

- LSTM in TensorFlow

Selectivity controls (**gates, 0 or 1**)

$$i_t = \sigma(W_i h_{t-1} + U_i x_t + b_i)$$

$$o_t = \sigma(W_o h_{t-1} + U_o x_t + b_o)$$

$$f_t = \sigma(W_f h_{t-1} + U_f x_t + b_f)$$

explicit remembering state

$$\tilde{c}_t = \phi(W h_{t-1} + U x_t + b)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

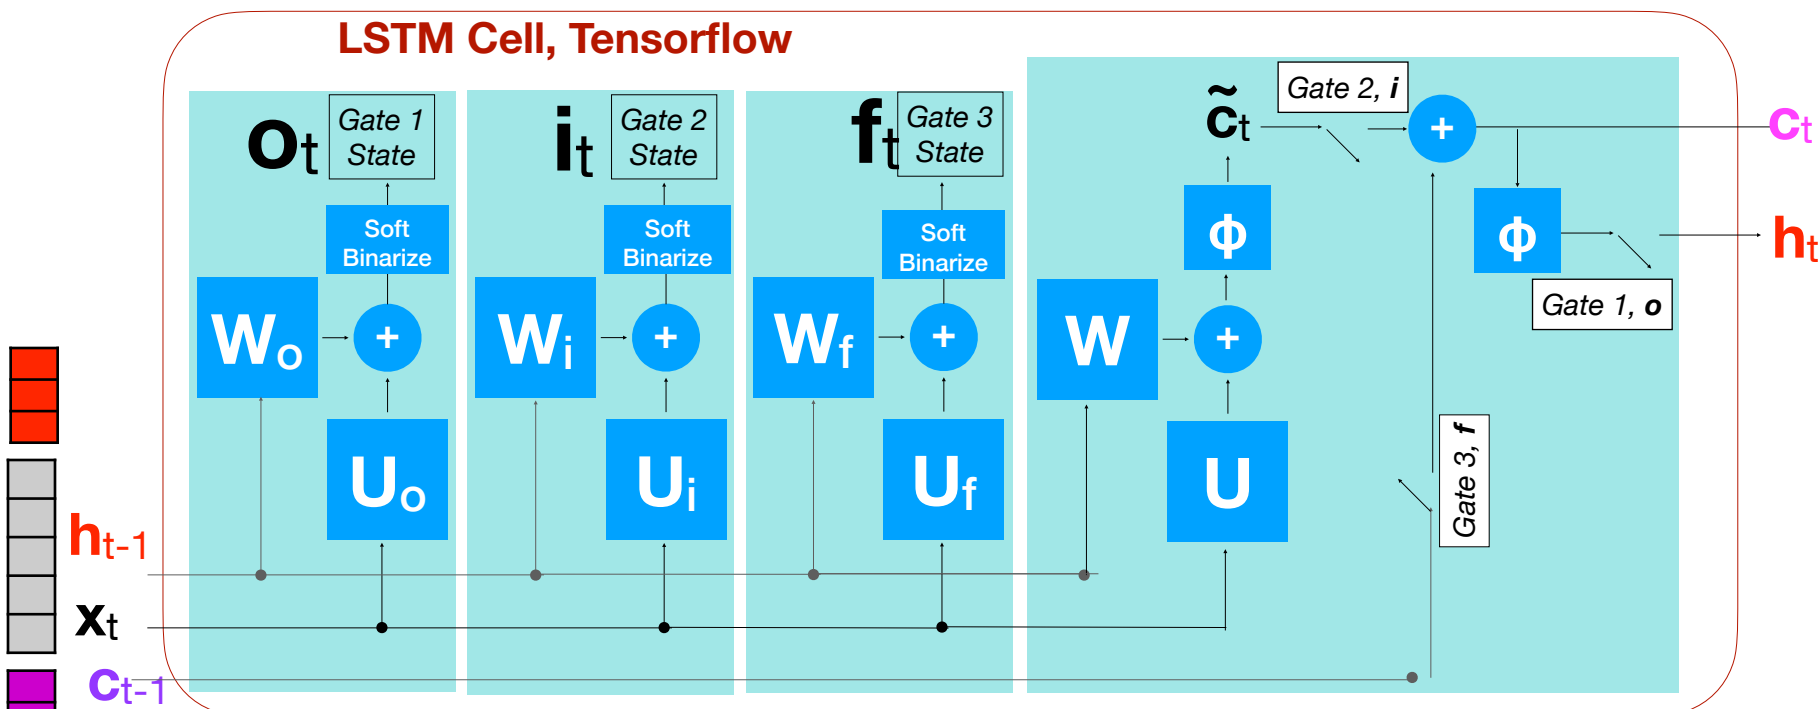
remember
previous state

update with
output, h_t

$$h_t = o_t \odot \phi(c_t)$$

get next h_t for
selecting gates

LSTM Cell, Tensorflow



LSTM Dropout

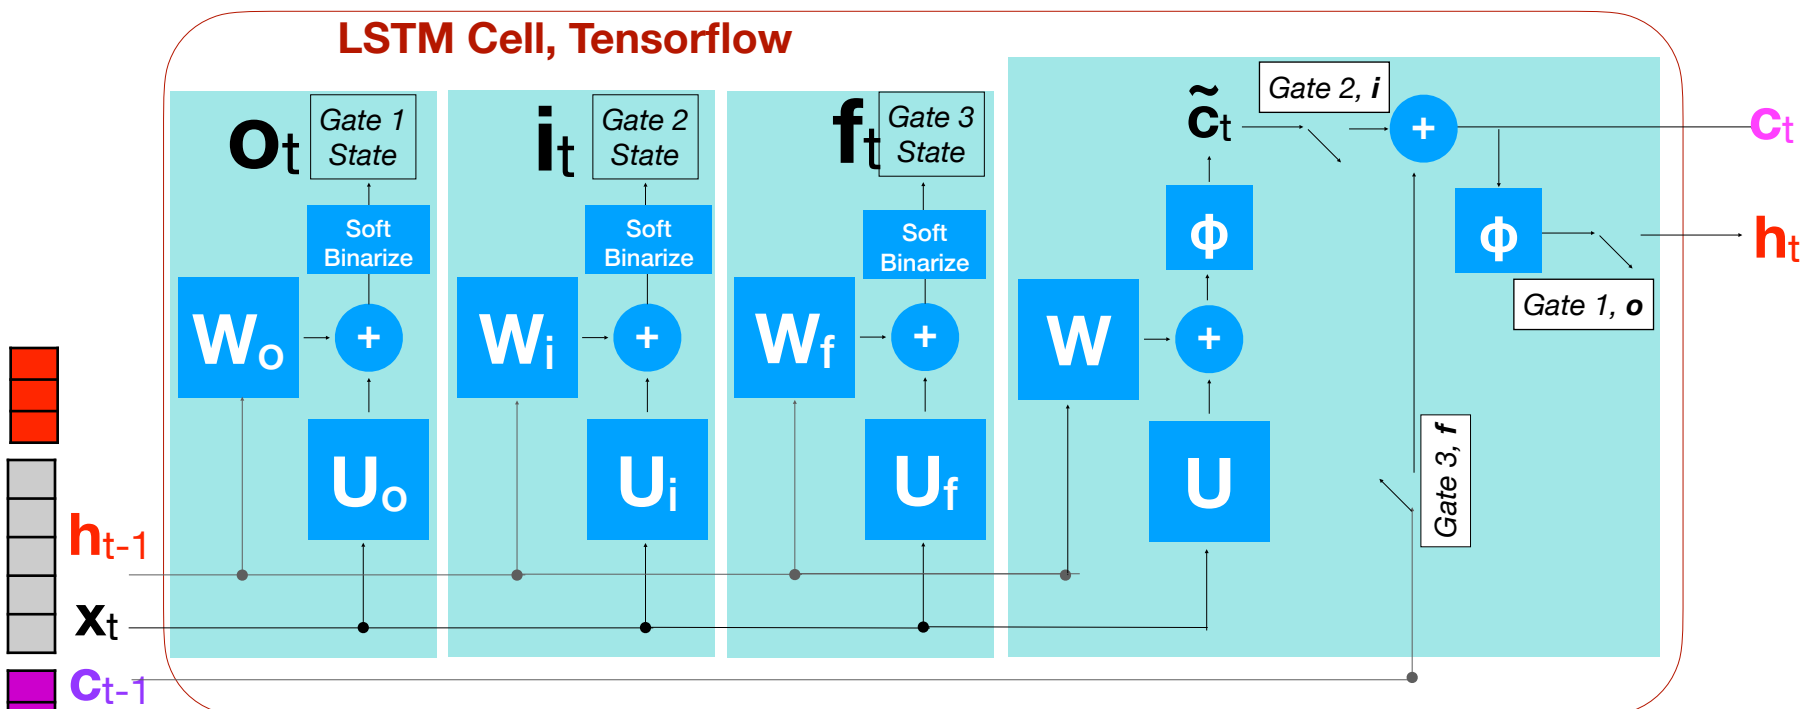
$$\begin{aligned}i_t &= \sigma(W_i h_{t-1} + U_i x_t + b_i) \\o_t &= \sigma(W_o h_{t-1} + U_o x_t + b_o) \\f_t &= \sigma(W_f h_{t-1} + U_f x_t + b_f)\end{aligned}$$

Recurrent
Dropout

Input
Dropout

The days of
training **without**
using **dropout** are
over.

LSTM Cell, Tensorflow



What to choose?

- There is no hard and fast rule
 - try both
 - basic LSTM has had great success
 - GRU also sometimes is easier to train
 - you will see many variations
 - peephole LSTM
 - hierarchical LSTM
 - and many more...

Many to one:
Simple RNNs
GRUs
LSTMs



More examples:


<https://github.com/tensorflow/tensorflow/tree/r0.11/tensorflow/examples/skflow>

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

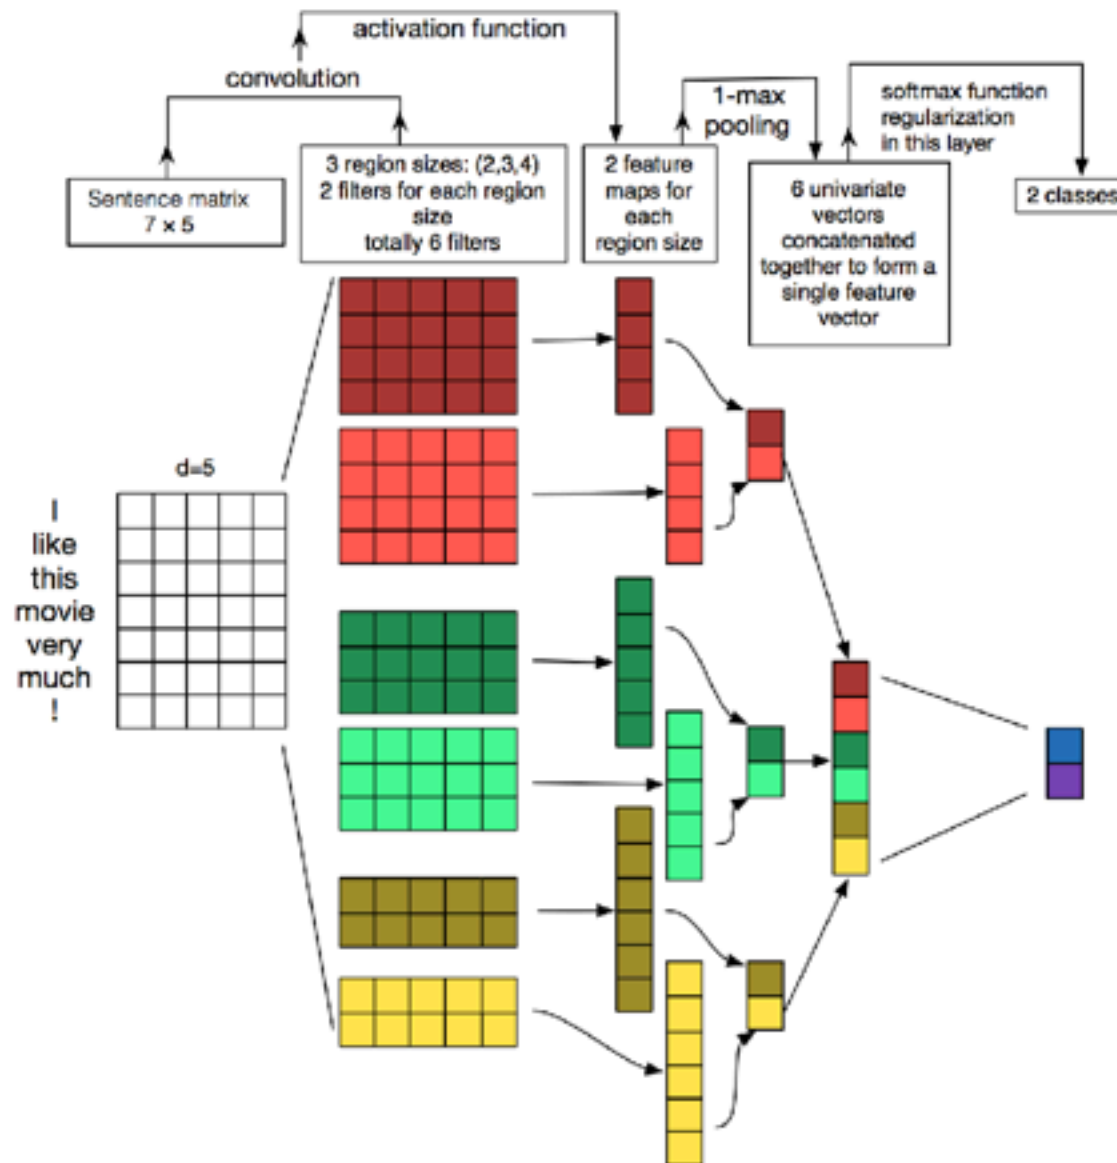
<http://machinelearningmastery.com/sequence-classification-lstm-recurrent-neural-networks-python-keras/>

Seq2Seq:

https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translation_word.py



CNNs with Multiple Region Sizes



Back to the CNN



More examples:

<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

Seq2Seq:

https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translation_word.py