Lecture Notes for

Machine Learning in Python

Professor Eric Larson

More Convolutional Networks

Class logistics and Agenda

- Wide/Deep Lab due soon!
- Agenda:
  - CNN Demo
  - History of CNNs
    - with Modern CNN Architectures
- Next Time:
  - More Advanced CNN Demo

Last Time:

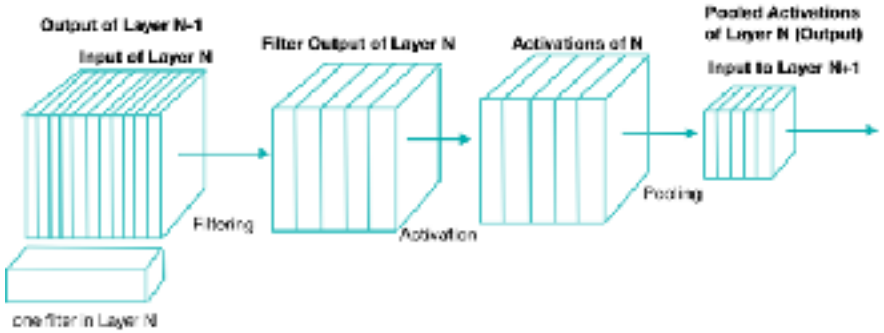

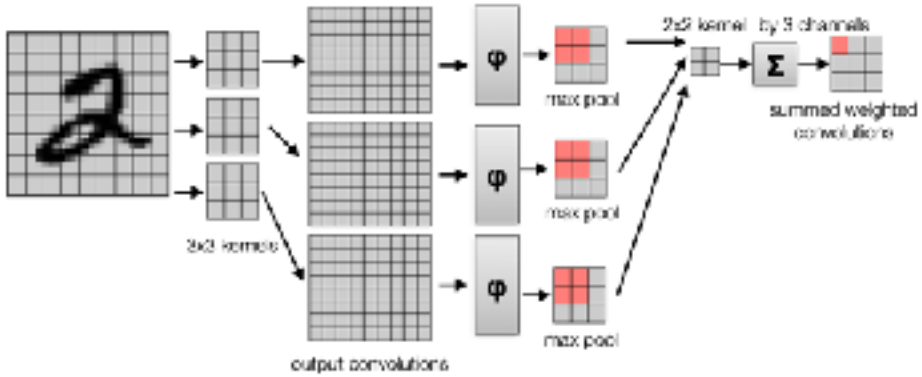

Structure of Each Tensor: Channels x Rows x Columns

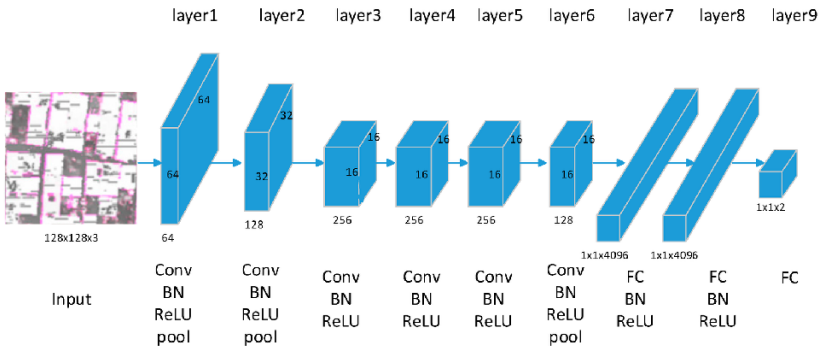

- Why perform pooling?
- Why max pooling?
  - reduce translation effects
  - param reduction

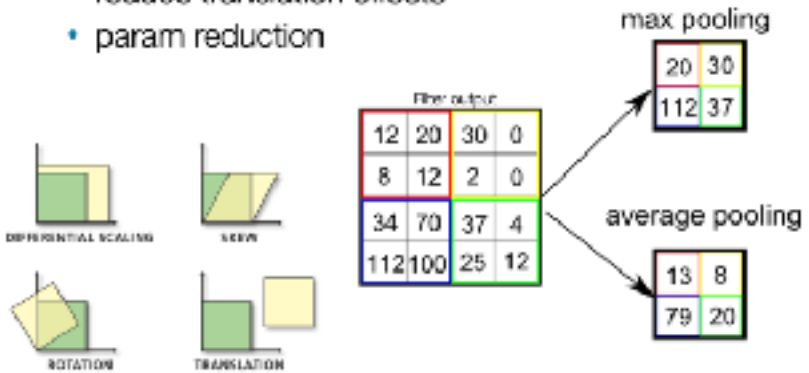

Demo

Convolutional Neural Networks
in TensorFlow
with Keras

`11. Convolutional Neural Networks.ipynb`

History of Convolutional Neural Networks

Thanks to machine-learning algorithms, the robot apocalypse was short-lived.

Types of CNN, 1988-1998

Heads Facebook AI Team

- **LeNet-1** (1988)
  - ~2600 params, not many layers
- **LeNet-5** (1998)
  - 7 layers, gets excellent MNIST performance
- Major contribution, general structure:
  - conv=>pool=>non-linearity=> …=>MLP

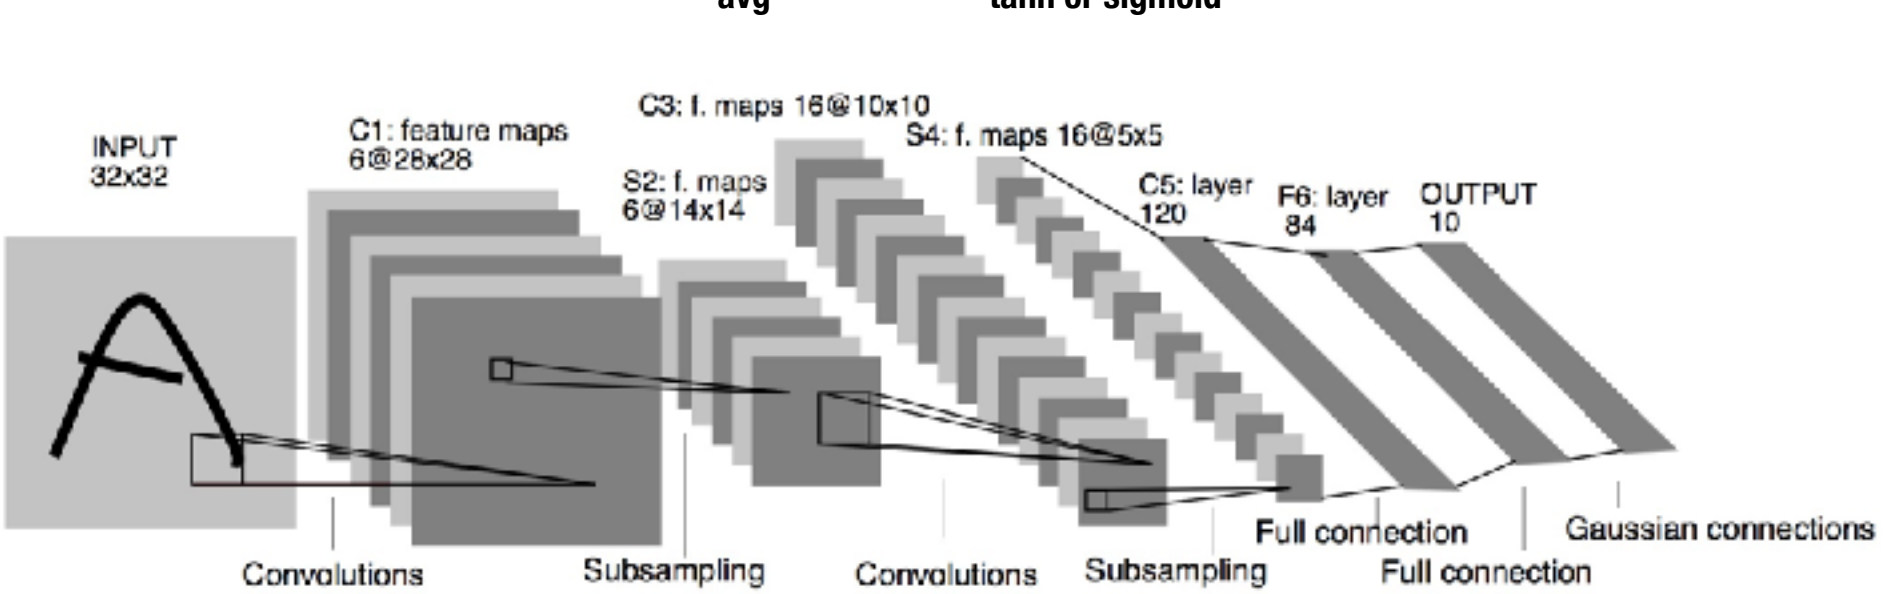

CNN History

- List of major breakthroughs from 1998 through 2010 in convolutional networks:

¯_(ツ)_/¯

  - 2010

Types of CNN, 2010

AI Researcher
IDSA, Switzerland

- **Ciresan Net**
- Publishes code for running CNN via GPU
  - Subsequently wins 5 international competitions
  - from stop signs => cancer detection
- Major contribution: NVIDIA parallelized training algorithms

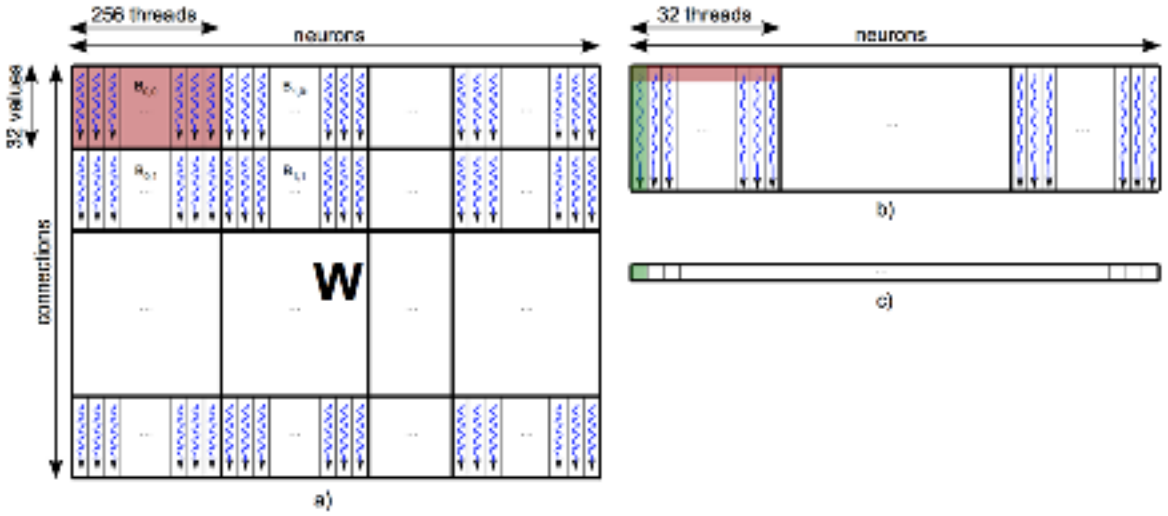

Figure 2: Forward propagation: a) mapping of kernel 1 grid onto the padded weight matrix; b) mapping the kernel 2 grid onto the partial dot products matrix; c) output of forward propagation.

https://arxiv.org/pdf/1003.0358.pdf

ImageNet Competition (2010-2016)

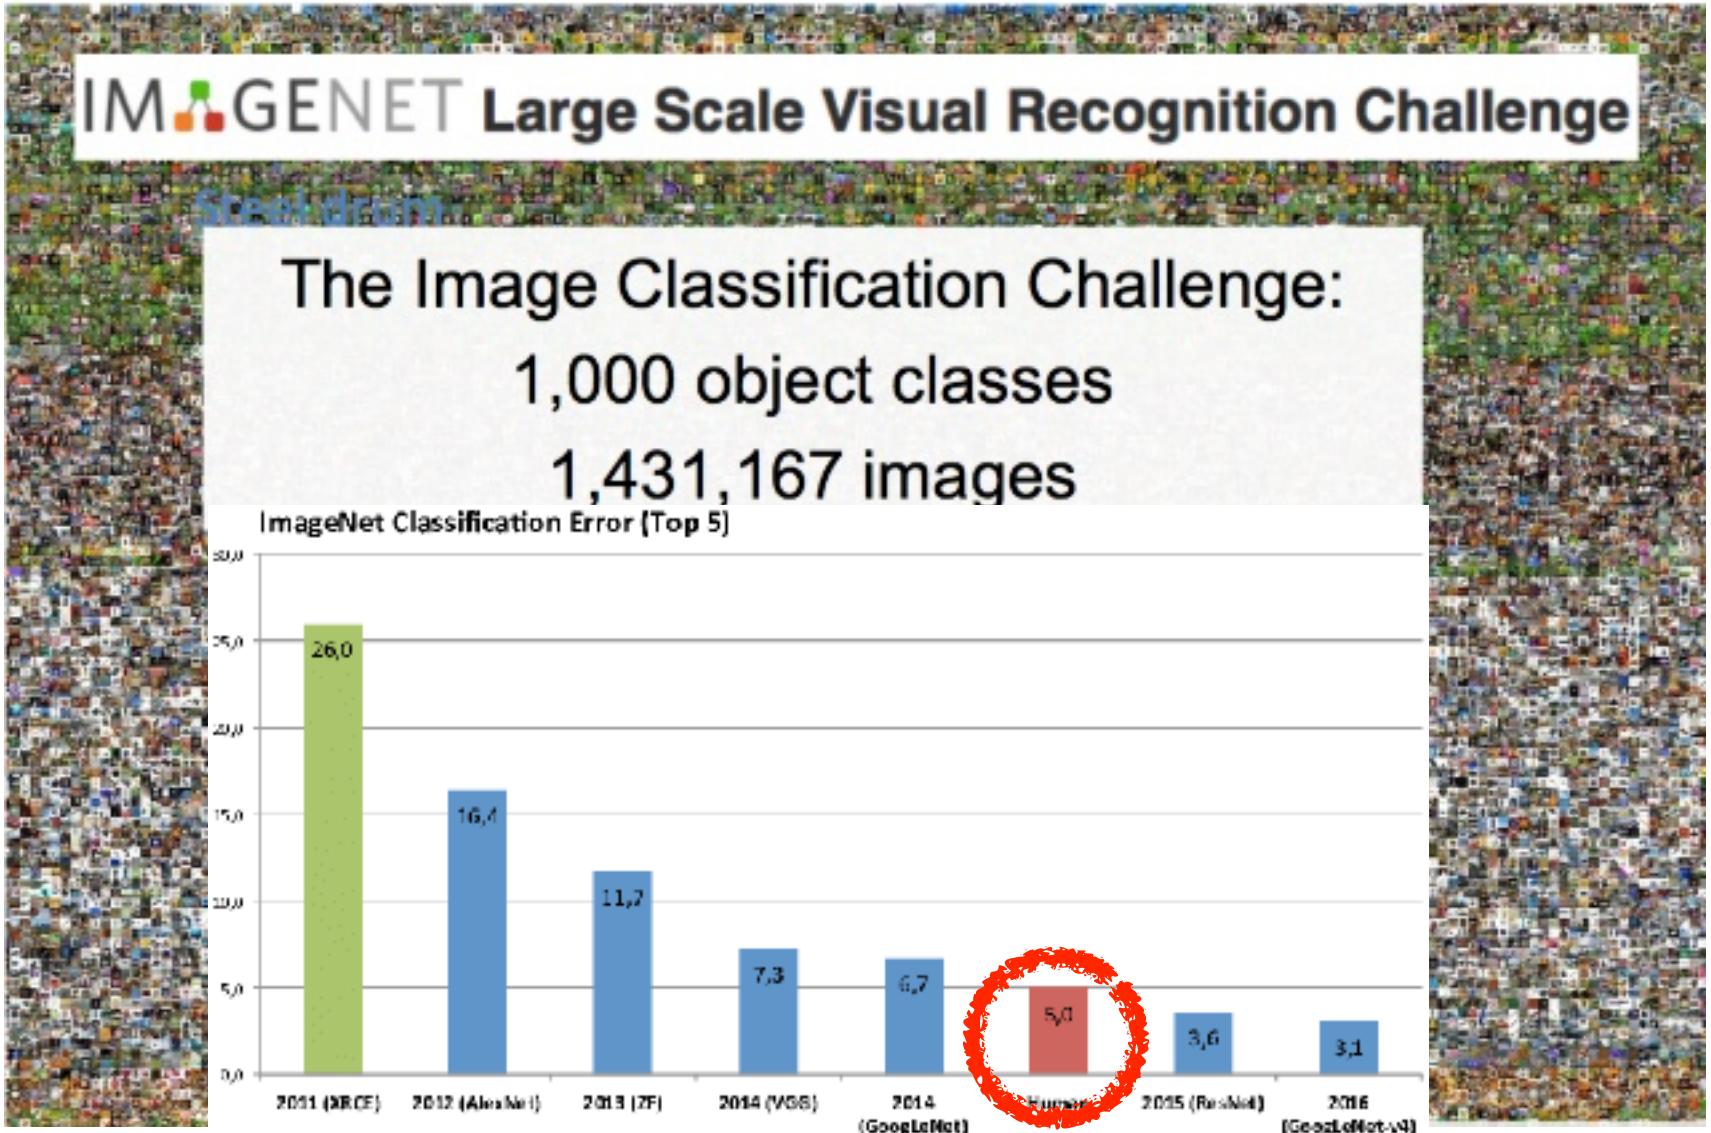

https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the_fig7_324476862

https://www.slideshare.net/nmhkahn/case-study-of-convolutional-neural-network-61556303

Types of CNN, 2012

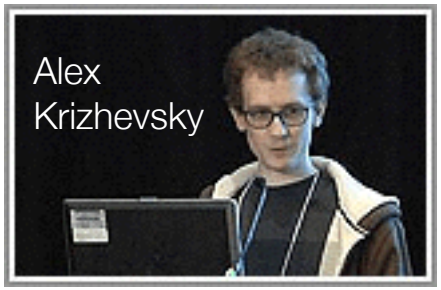

Google

- **AlexNet**, Hinton is mentor
  - wins ImageNet competition
- Major contributions:
  - dropout for regularization
  - systematic use of ReLU
  - data expansion
  - ***overlapping max pool***

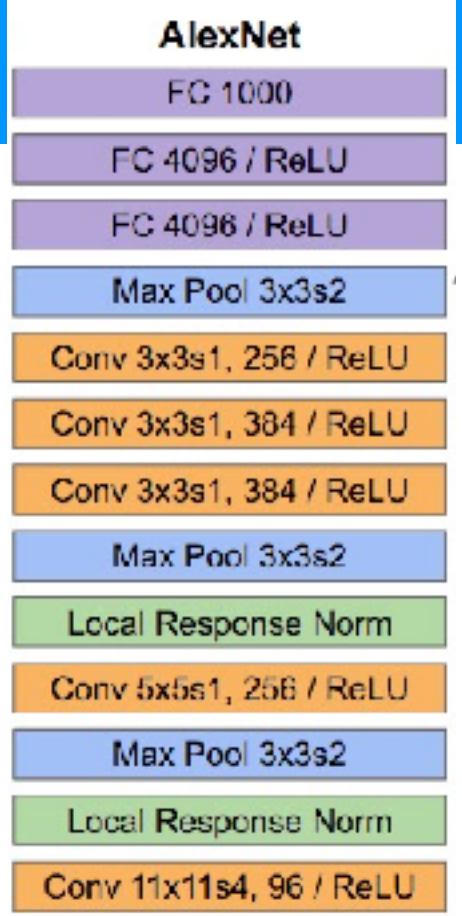

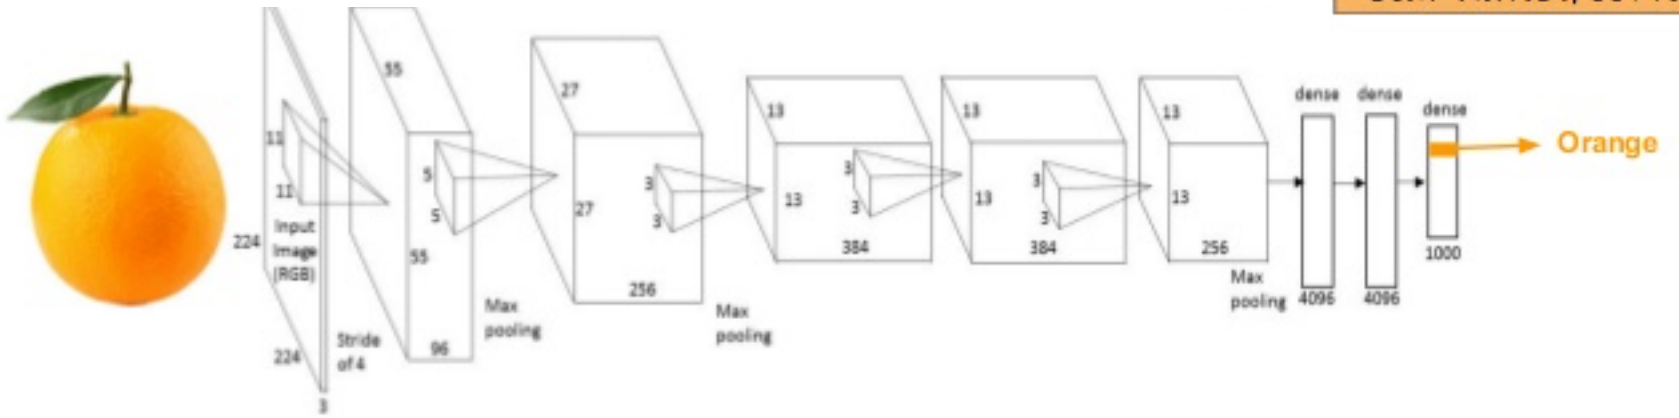

https://culurciello.github.io/tech/2016/06/04/nets.html

Warning


DON'T TRY WHAT YOU'RE ABOUT TO SEE AT HOME
WE'RE WHAT YOU CALL EXPERTS
WeKnowMemes

Types of CNN, 2013

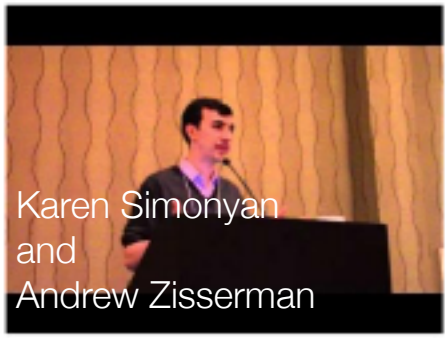

UNIVERSITY OF OXFORD

Google DeepMind

- Oxford **VGG Net** (Visual Geometry Group)
- Major contributions:
  - small cascaded kernels
  - way more layers (19 versus ~7)
  - "emulates" biology "better"
  - trained on NVIDIA GPUs for 2-3 weeks

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 **LRN**	conv3-64 **conv3-64**	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 **conv3-128**	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 **conv1-256**	conv3-256 conv3-256 **conv3-256**	conv3-256 conv3-256 conv3-256 **conv3-256**
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 **conv1-512**	conv3-512 conv3-512 **conv3-512**	conv3-512 conv3-512 conv3-512 **conv3-512**
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 **conv1-512**	conv3-512 conv3-512 **conv3-512**	conv3-512 conv3-512 conv3-512 **conv3-512**
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: **Number of parameters** (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

http://www.robots.ox.ac.uk/~vgg/practicals/cnn/

Types of CNN, 2014

- Network in Network **NiN**
  - or MLPConv

https://arxiv.org/pdf/1312.4400.pdf

Types of CNN, 2014

- Network in Network

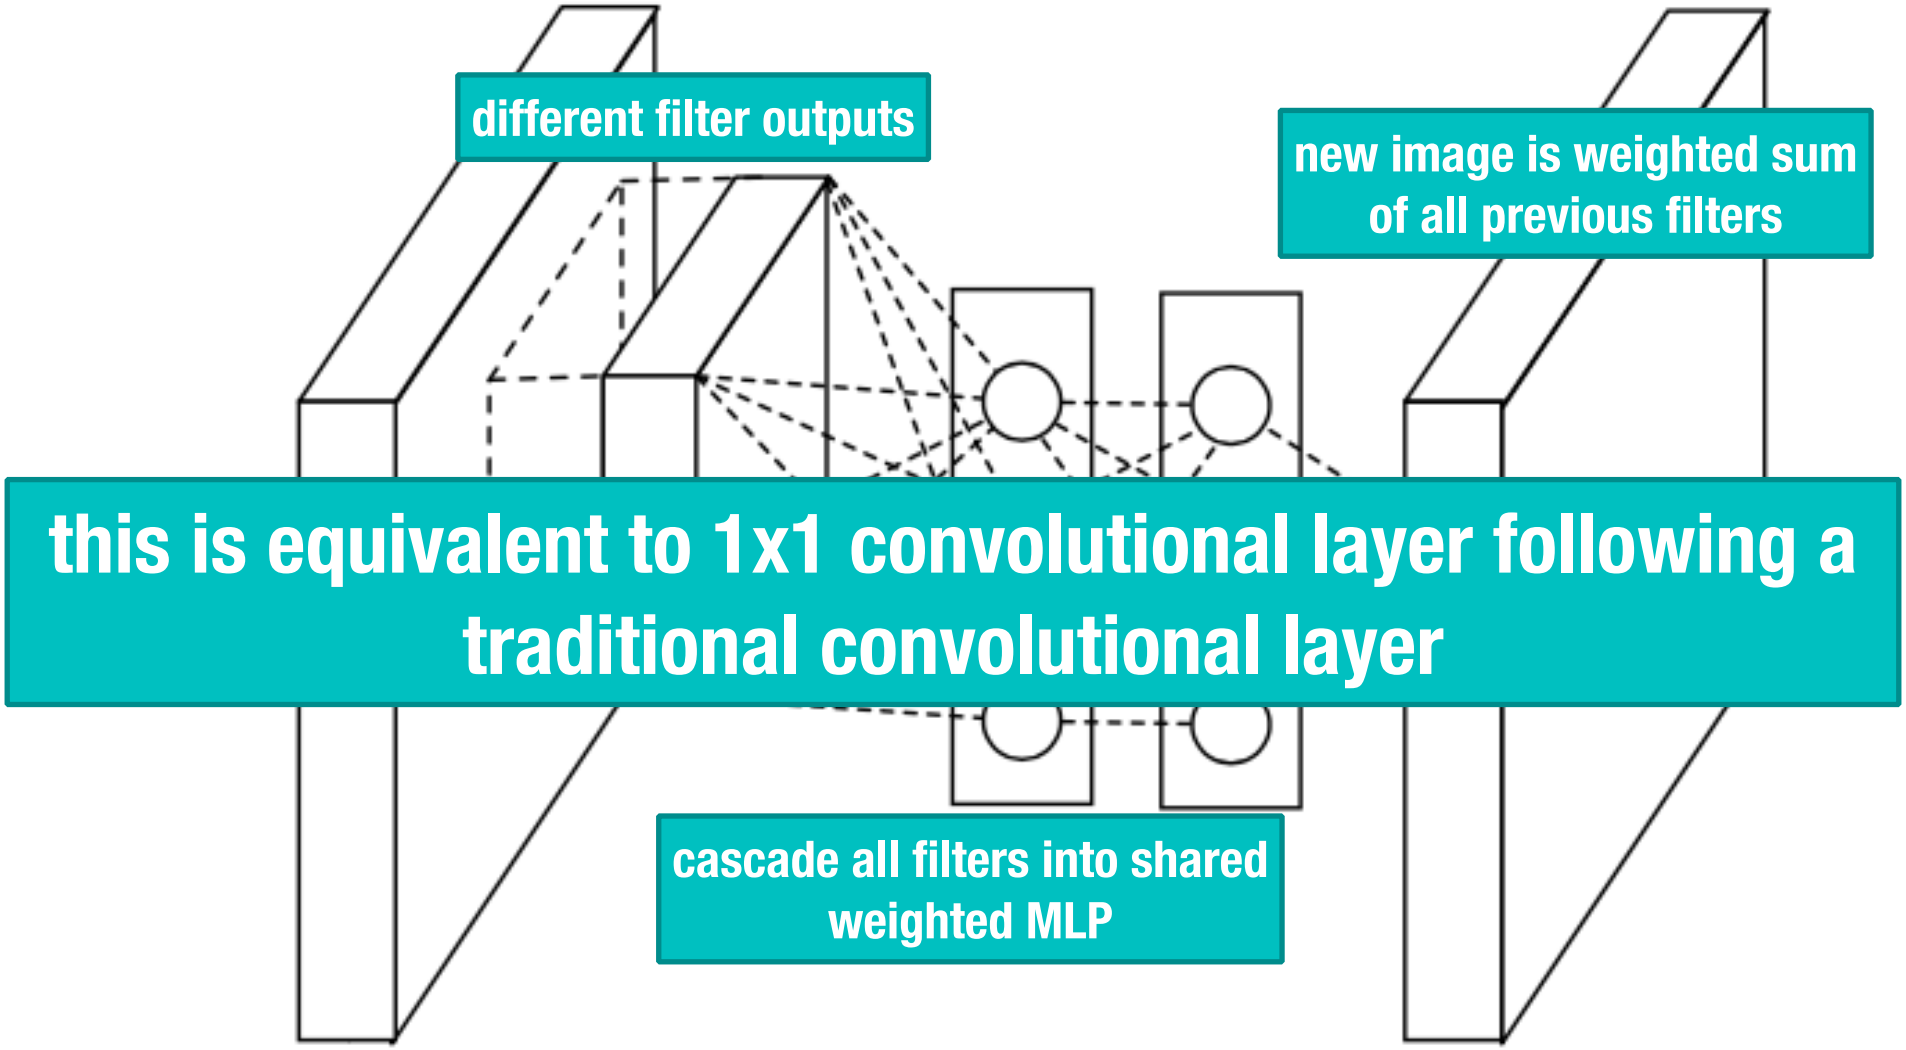

https://arxiv.org/pdf/1312.4400.pdf

NiN, expanded view

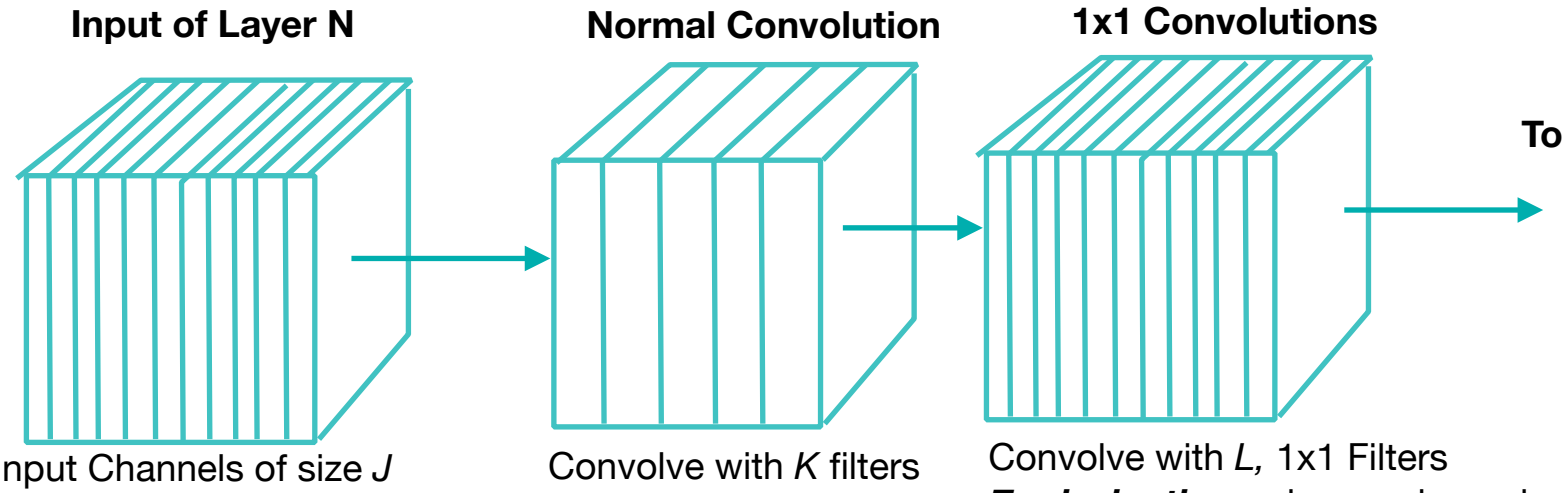

Structure of Each Tensor: Channels x Rows x Columns

Types of CNN, 2014

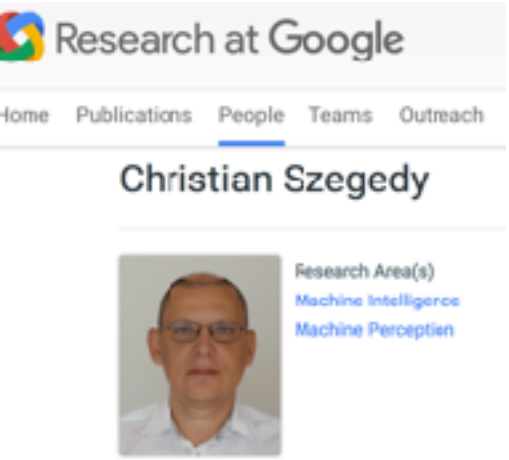

- **GoogLeNet**
  - or **Inception V1**
- Major contribution:
  - bottleneck layering
  - parallel NiN

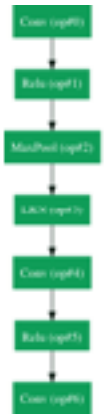

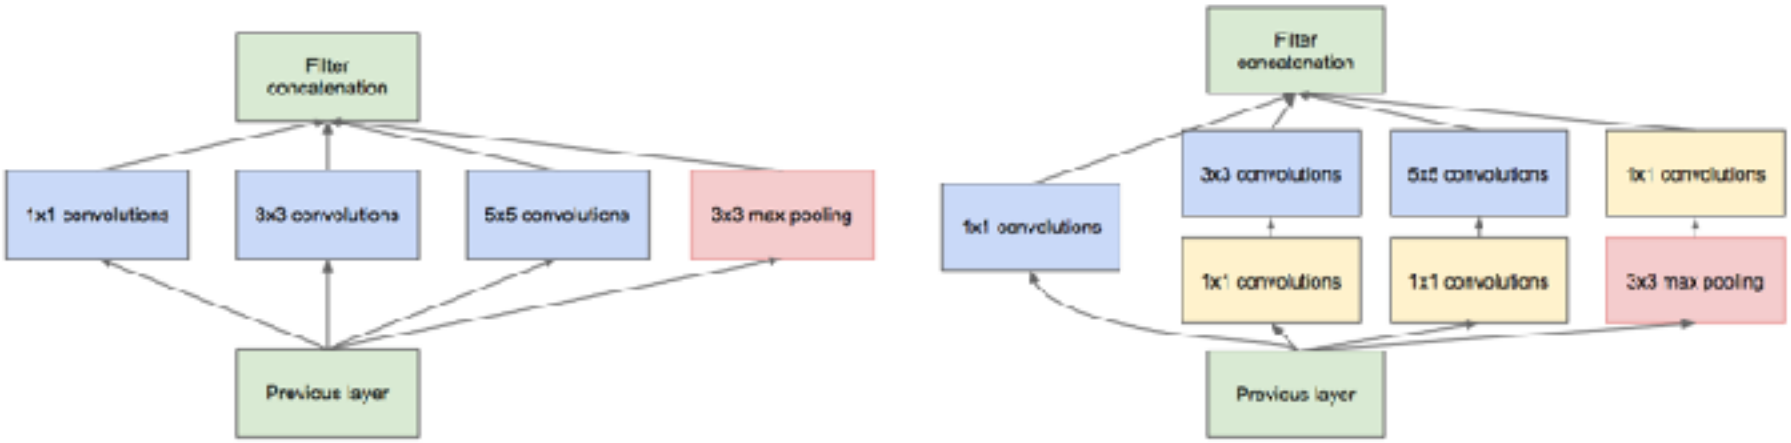

(a) Inception module, naïve version

(b) Inception module with dimension reductions

Figure 2: Inception module

https://arxiv.org/pdf/1409.4842.pdf

Types of CNN, 2015 February and December

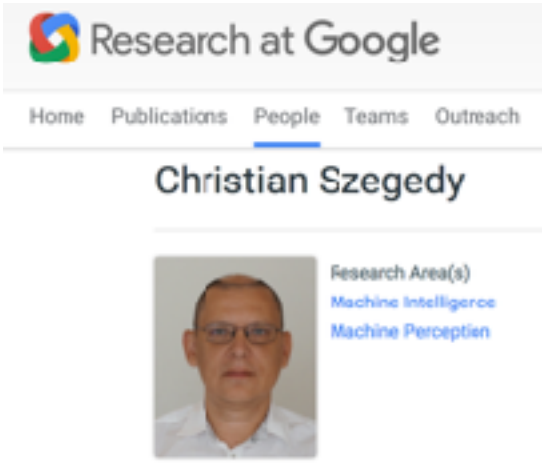

- **Inception V2**, Inception V1 with batch normalization
- **Inception V3**:
  - replace 5x5 with multiple 3x3

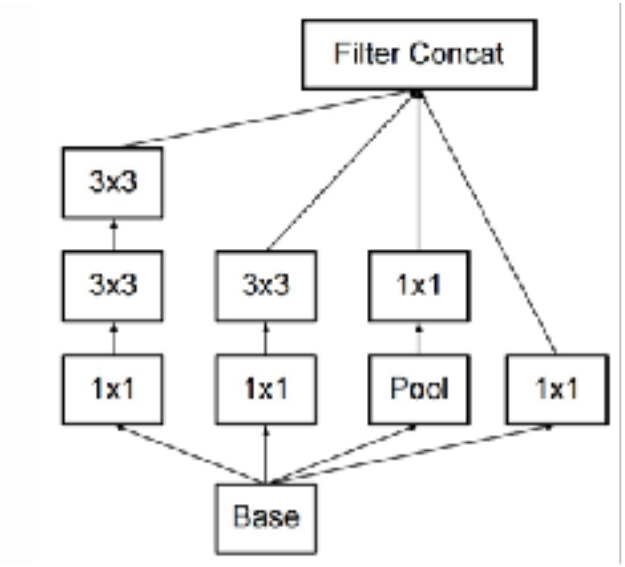

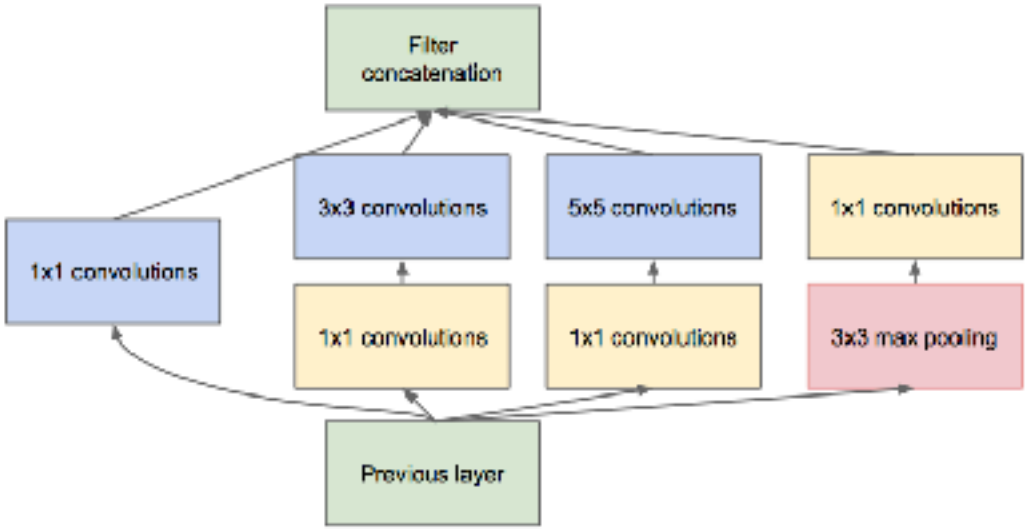

https://arxiv.org/pdf/1512.00567.pdf

Types of CNN, 2015 December

Microsoft Research

- Major Contributions:
  - "ensembles" not strictly sequential
  - "bio-plausible" with feedback

- **ResNet**
  - Parametric ReLU
    PReLU: adaptive trained slope

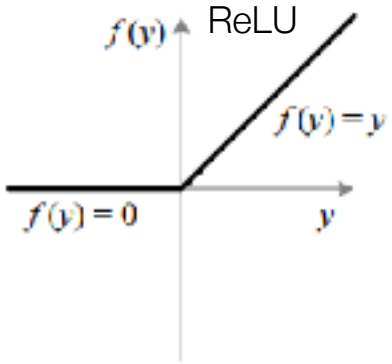

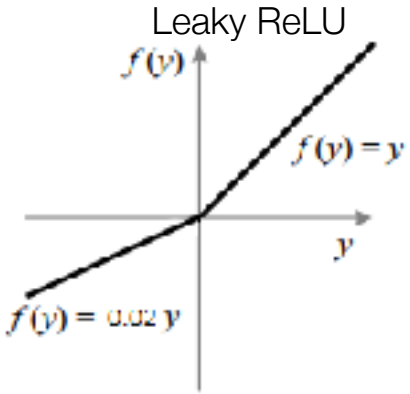

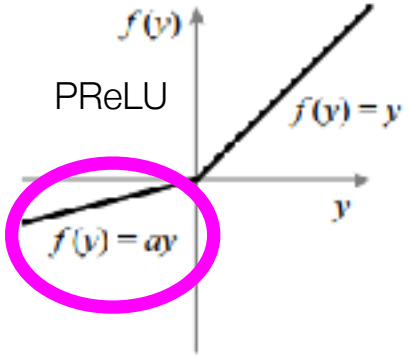

- NiN: triple bypass layer
  - similar to bottelneck

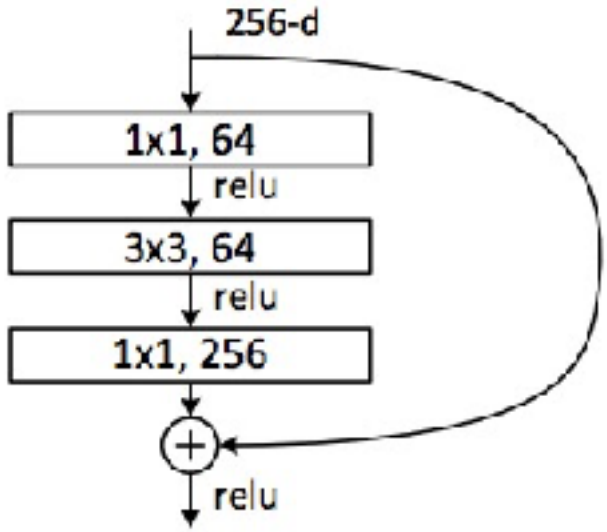

Residual Connection, expanded view

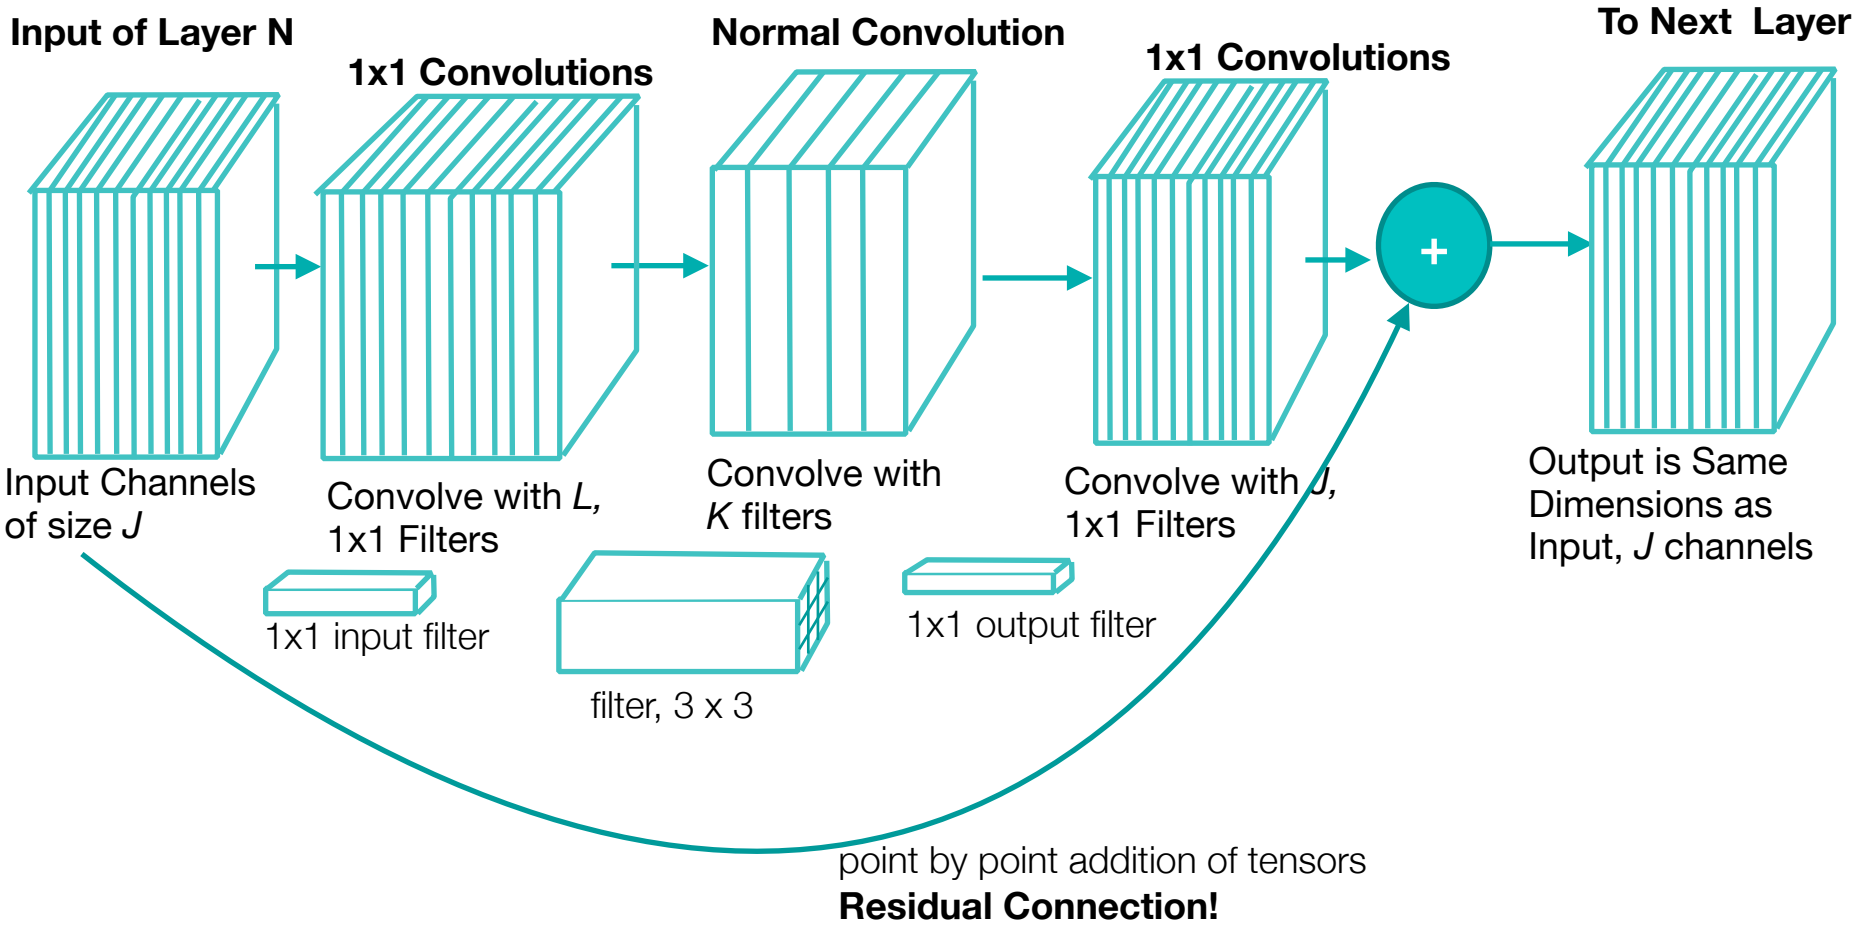

Back Propagation: Two paths, including one without ANY operations that cause the gradient to vanish…

Types of CNN, 2017 April

Xception

- Major Contributions:
  - combining branching / residual blocks
  - separable convolutions

Francois Chollet
Google

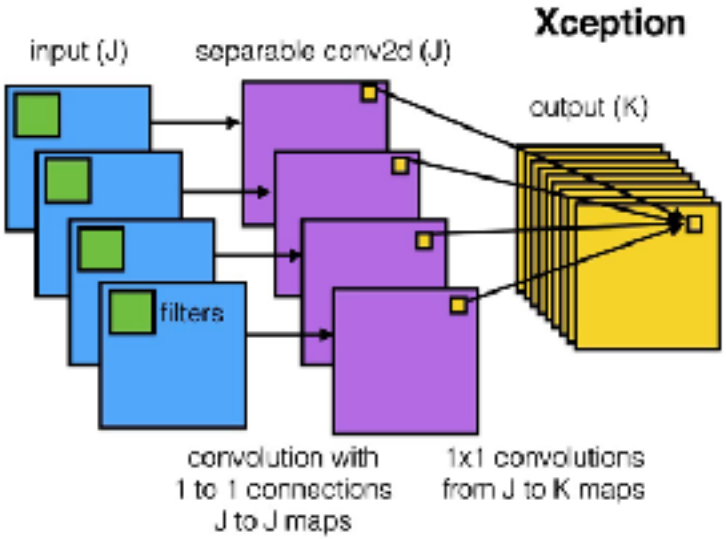

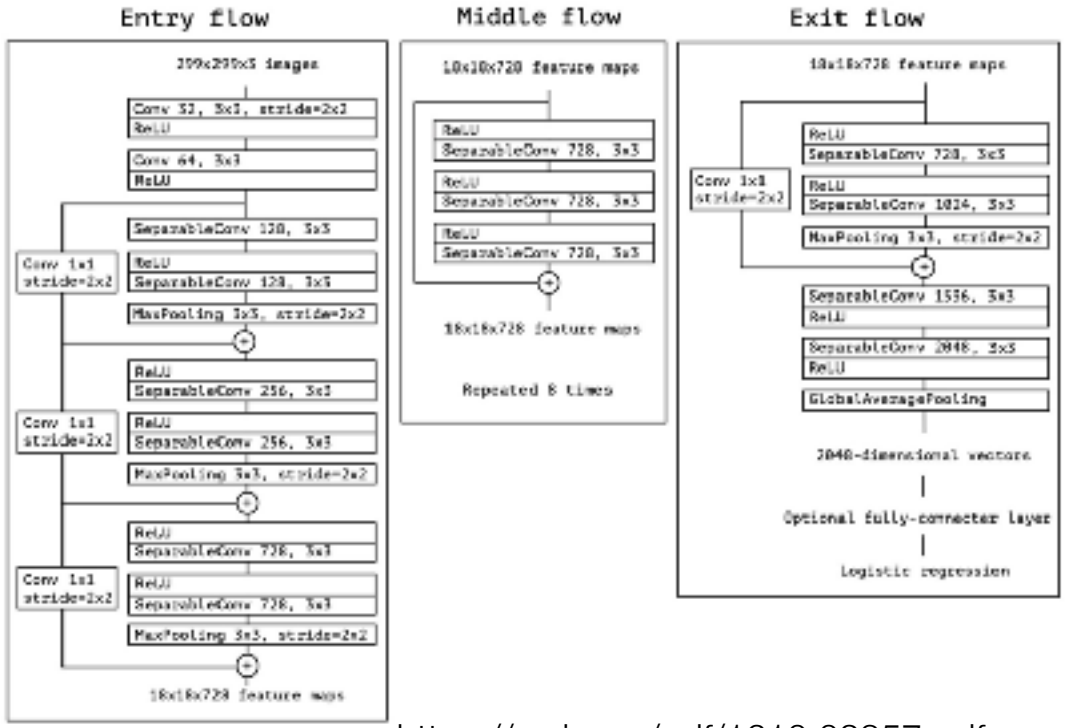

https://arxiv.org/pdf/1610.02357.pdf

Separable Convolution Primer

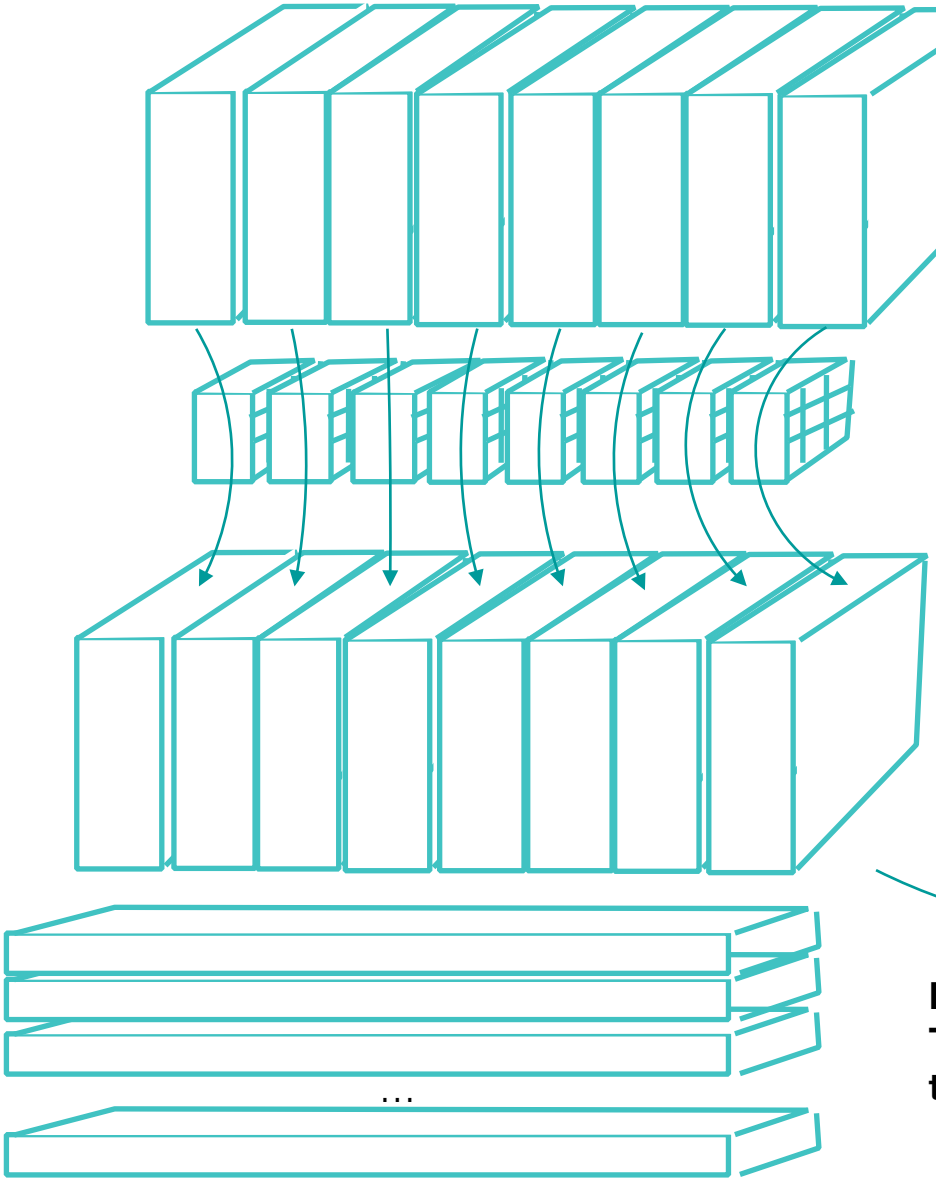

How big are these networks?

How big are these networks?

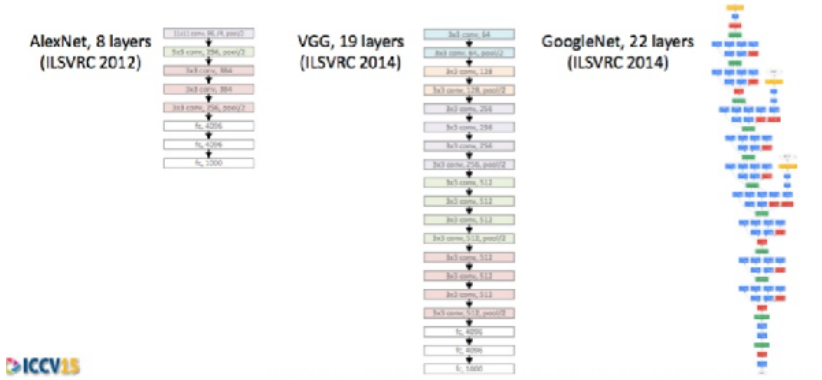

ResNet, 152 layers
(ILSVRC 2015)

34-layer residual

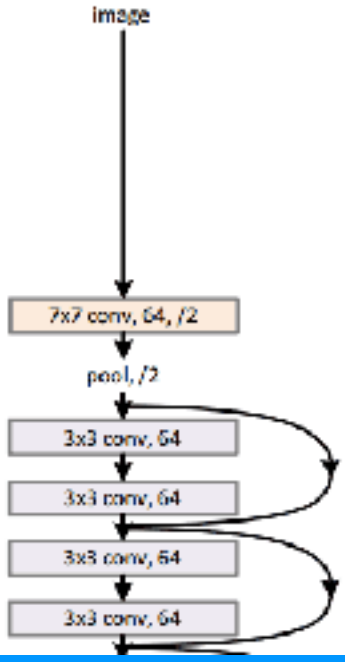

Self Test

- We have seen a lot of different networks.
- The most important concept to understand in using convolutional neural networks is:
  - A. Use proper initialization of layers
  - B. Have plenty of data or use expansion
  - C. Set aside time for training
  - D. Use batch normalization

Next Time:

- CNN Lab Discussion (Town Hall)

- Intro to Recurrent Neural Network Architectures
  - RNNs, GRUs, LSTMs
  - Training for characters